

India Launch: Detailed Step-by-Step Plan

(All acronyms expanded and each step elaborated with “what” and “why”)

1. Prerequisites

1. Amazon Web Services (AWS) Account & Identity and Access Management (IAM) Permissions

- **What:**
 - An AWS account with an IAM user or role that has permissions to create and manage:
 - **App Runner** (fully managed container service)
 - **Amazon Relational Database Service (RDS)**
 - **Virtual Private Cloud (VPC)** and related networking
 - **Amazon Route 53** (DNS service)
 - **AWS Secrets Manager** (credential storage)
- **Why:**
 - Ensures you have the “right” (least-privilege) access to provision, configure, and secure all necessary cloud resources without using overly permissive credentials.

2. Domain Hosted in Amazon Route 53

- **What:**
 - A registered domain (for example, `yourdomain.com`) with a **hosted zone** configured in Amazon Route 53.
 - You’ll use a subdomain like `erp.yourdomain.com` for your ERP’s public endpoint.
- **Why:**
 - Allows you to map a friendly, branded, and secure URL (with HTTPS) to your App Runner service later.

3. AWS Command Line Interface (CLI) (Optional but Recommended)

- **What:**
 - The AWS CLI installed on your local machine or available via AWS CloudShell.
 - Configured with your IAM user’s access key and secret key.
 - **Why:**
 - Facilitates scripting and automation of repetitive tasks (e.g., creating VPCs, RDS instances) and enables infrastructure-as-code workflows.
-

2. Network & Security Setup

1. Virtual Private Cloud (VPC) Selection or Creation

- **What:**
 - Use the **default VPC** in the **Asia Pacific (Mumbai)** Region (`ap-south-1`) or create a new VPC with at least two **subnets** spread across different **Availability Zones (AZs)**.
- **Why:**
 - A VPC provides an isolated network environment.

- Multiple AZs ensure high availability—if one AZ fails, resources in the other AZ can continue serving traffic.
 - 2. **Security Group for RDS (Database) Access**
 - **What:**
 - Create a **security group** (e.g., named `erp-rds-sg`) within your VPC.
 - Configure **inbound rules** to allow only PostgreSQL traffic (TCP port 5432) from the **App Runner VPC Connector**'s CIDR block (the IP range App Runner will use).
 - Keep **outbound rules** open (default “all traffic allowed”) so the database can communicate for backups, monitoring, etc.
 - **Why:**
 - Security groups act as virtual firewalls.
 - Restricting inbound access to only your application prevents unauthorized database connections.
 - 3. **(Optional) VPC Connector for App Runner**
 - **What:**
 - Define an **App Runner VPC Connector** resource that integrates App Runner services into your private VPC network.
 - **Why:**
 - By default, App Runner services run in a managed network.
 - Attaching a VPC Connector gives your containerized application private network access to RDS—no need to expose the database publicly.
-

3. Provision Amazon RDS PostgreSQL (Multi-AZ)

- **What:**
 - In the RDS console, choose “**Create database**” with engine **PostgreSQL**.
 - Select the **db.t3.medium** instance class (2 vCPU, 4 GiB RAM) with **Multi-AZ deployment** enabled in the Mumbai Region (`ap-south-1`).
 - Configure storage: 20 GiB GP3 (General Purpose SSD) with auto-scaling.
 - Set **backup retention** to 7 days and enable **performance insights** if desired.
 - **Why:**
 - **Multi-AZ** ensures automatic failover to a standby instance in another AZ if the primary fails—critical for high availability.
 - Daily automated backups let you restore point-in-time.
 - Performance insights help you monitor and tune database load.
-

4. Store Database Credentials Securely

- **What:**
 - In **AWS Secrets Manager**, create a new secret of type “**Credentials for RDS database**”.
 - Enter the master **username**, **password**, and the RDS **endpoint** (hostname).
 - Name it something like `erp/production/db`.
- **Why:**

- Secrets Manager encrypts credentials at rest and manages rotation policies.
 - Your application can retrieve them at runtime via IAM-based requests—no hard-coding of sensitive information in code or environment variables.
-

5. Deploy Your Application to AWS App Runner

- **What:**
 1. In the **App Runner** console, click “**Create service.**”
 2. Select your **source**:
 - **GitHub repository** (connected via OAuth) with build instructions (e.g. Dockerfile), or
 - **Amazon Elastic Container Registry (ECR)** image URI.
 3. Name the service `erp-frontend-india-launch` and set the **Region** to `Mumbai (ap-south-1)`.
 4. Configure **instance size** to **1 vCPU / 2 GiB RAM** and **auto-scaling** (min 1, max 2).
 5. Under **Networking**, attach your **VPC Connector** so the service can talk privately to RDS.
 6. Define **environment variables** (pointing to Secrets Manager ARNs or using parameter names):
 - `DB_HOST` → RDS endpoint
 - `DB_NAME`, `DB_USER`, `DB_PASS` → retrieved at runtime from Secrets Manager
 7. Enable **logs** to CloudWatch and accept default HTTP **health check** settings.
 - **Why:**
 - App Runner is a fully managed platform—no servers to patch or scale manually.
 - Built-in auto-scaling and health checks keep your service resilient.
 - Private networking with the VPC Connector secures database traffic.
-

6. Point Your Custom Domain & Enable HTTPS

1. **Add Custom Domain in App Runner**
 - In the App Runner service details, under “**Custom domains,**” add `erp.yourdomain.com`.
 - App Runner automatically provisions an **SSL/TLS certificate** via AWS Certificate Manager.
 2. **Create Alias Record in Route 53**
 - In your Route 53 hosted zone, create an **A – IPv4 address** record:
 - **Name:** `erp`
 - **Alias:** Yes → select the App Runner custom domain endpoint.
- **Why:**
 - Provides a user-friendly URL.
 - Automatic certificate provisioning ensures all traffic is encrypted (HTTPS) without manual cert management.

7. Validation & Smoke Test

- **What:**
 1. Visit `https://erp.yourdomain.com` in your browser—confirm the app loads.
 2. Through the UI, execute key flows such as:
 - **Create Company** → new tenant setup.
 - **List Accounts** → fetch Chart of Accounts from the database.
 3. Review **CloudWatch Logs** for any errors or timeouts.
- **Why:**
 - Verifies end-to-end connectivity (App Runner → RDS).
 - Catches misconfigurations before handing over to users.

8. Monitoring, Backups & Security Hardening

1. **Amazon RDS Monitoring & Backups**
 - Ensure **automated backups** are enabled (7-day retention).
 - Turn on **Enhanced Monitoring** at 5-second granularity.
 - Create **CloudWatch Alarms** for:
 - CPU utilization > 80%
 - Storage used > 75% of allocated
 - Replica lag (if future replica exists) > 1 second
 2. **App Runner Metrics**
 - Configure CloudWatch Alarms on:
 - Request count anomalies
 - High error rates (5xx status codes)
 - Latency spikes
 3. **IAM & Secrets Manager Practices**
 - **Rotate** the database secret every 30–90 days.
 - Limit the App Runner service role to only:
 - `secretsmanager:GetSecretValue` for the specific secret ARN
 - RDS connect permissions (via security group rules)
- **Why:**
 - Early-warning on resource exhaustion or performance degradation.
 - Least-privilege IAM reduces blast radius if credentials or roles are compromised.

9. Next Steps & Future Evolution

1. **Continuous Integration / Continuous Deployment (CI/CD)**
 - Connect your Git repository to App Runner so that pushes to `main` trigger new builds and automatic rollouts.

2. Staging & Production Environments

- Replicate this setup in separate AWS accounts or VPCs.
- Use parameterized CloudFormation or Terraform stacks with distinct environment variables.

3. Asia Expansion (Read-Replica & Edge Routing)

- When you need low-latency reads in Southeast Asia:
 - Repeat RDS provisioning as a **cross-Region read replica** in Singapore (ap-east-1).
 - Deploy a second App Runner service in that region, pointing at the replica endpoint.
 - Enable **AWS Global Accelerator** for intelligent routing to the nearest healthy endpoint.

- **Why:**

- Phased rollout minimizes initial cost and complexity.
- Allows you to validate India-only operations before scaling geographically.
- Ensures global users see sub-200 ms response times as you expand.